

Insurance Enrollment Prediction - Technical Report

1. Introduction

This project aims to predict whether an employee will opt in to a new voluntary insurance product based on demographic and employment-related features. We build a machine learning pipeline covering data preprocessing, model training, evaluation, and deployment via a REST API.

2. Dataset Overview

- **Source:** Synthetic employee census-style dataset
- **Size:** ~10,000 rows, representing employee records
- **Features:**
 - `employee_id` (unique identifier)
 - `age` (numeric)
 - `gender` (categorical)
 - `marital_status` (categorical)
 - `salary` (numeric)
 - `employment_type` (categorical)
 - `region` (categorical)
 - `has_dependents` (binary encoded as categorical)

- `tenure_years` (numeric)
 - `enrolled` (target: 0 or 1)
-

3. Data Processing

- Handled missing values and ensured consistent data types.
 - Encoded categorical variables using one-hot encoding for model compatibility.
 - Scaled numerical features using `StandardScaler` to normalize data.
 - Performed train-test split maintaining representative data distributions.
 - Saved preprocessing pipeline using `joblib` for reuse during inference.
-

4. Model Development

- Chose **XGBoost** as the primary model due to its proven performance on tabular data and ability to handle mixed feature types.
 - Trained the model on processed training data.
 - Hyperparameter tuning was performed via MLflow to optimize model parameters and track experiments.
-

5. Model Evaluation

- Evaluation metrics computed on test set:
 - Accuracy

- Precision
 - Recall
 - F1-score
 - ROC AUC
 - Visualized results with interactive Plotly charts:
 - Bar chart of evaluation metrics
 - Confusion matrix heatmap
 - ROC curve
 - Saved all plots in both HTML and PNG formats for reporting.
-

6. Model Persistence

- Saved trained model and preprocessing pipeline as pickle files (`model.pkl` and `preprocessor.pkl`) for reproducibility.
 - Enabled easy model loading during serving to avoid retraining.
-

7. Deployment via REST API

- Developed a **FastAPI** REST endpoint `/predict` to serve real-time predictions.
- API accepts employee features as JSON, applies preprocessing, and returns enrollment probability.
- Tested with sample `curl` commands.
- Input schema validation implemented using Pydantic models for robust request handling.

8. Experiment Tracking

- Integrated **MLflow** for:
 - Logging hyperparameters
 - Recording evaluation metrics
 - Storing model artifacts
- Enables reproducibility and comparison across multiple training runs.

9. Key Takeaways

- The XGBoost model effectively predicts voluntary insurance enrollment using demographic and employment data.
- Proper data preprocessing and feature encoding are crucial for model performance.
- Interactive visualizations aid in interpreting model strengths and weaknesses.
- REST API deployment enables practical integration of ML predictions into business workflows.
- MLflow experiment tracking facilitates systematic hyperparameter tuning and model management.

10. What to Do Next with More Time

- **Data Enrichment:** Incorporate additional employee features like performance ratings, historical claims, or behavioral data for better predictions.
- **Feature Engineering:** Use domain knowledge to create composite features or embeddings to capture complex relationships.

- **Advanced Models:** Experiment with deep learning models or ensemble stacking to further improve accuracy.
 - **Automated Hyperparameter Optimization:** Use Bayesian optimization frameworks (e.g., Optuna) for more efficient tuning.
 - **API Enhancements:** Add authentication, rate limiting, and logging for production readiness.
 - **CI/CD Pipeline:** Automate testing, deployment, and monitoring of the model with continuous integration/continuous deployment tools.
 - **Scalability:** Containerize the API using Docker and deploy on cloud platforms with autoscaling.
 - **User Interface:** Build a simple frontend to interact with the API and visualize predictions.
 - **Explainability:** Integrate SHAP or LIME for model interpretability to explain individual predictions to stakeholders.
-

Appendix

- **Code repository:** <https://github.com/CSKacas/Insurance-Enrollment-Prediction.git>
- **Data file:** `employee_data.csv` (synthetic dataset)
- **Environment:** Python 3.9, XGBoost, FastAPI, MLflow, Plotly